Bolted System: Auto-deployment Cloud Project
(Sprint 5)

Vidya Anandamurali
Pei Jia
Yuxi Jiang
Jiangnan Zou

Project Description (Recap)

Automate the deployment of Bolted which consists of:

- Container image of each component of the Bolted system (HIL, BMI, Keylime and orchestration)
- Automated deployment of component containers on a cloud platform

BU

Project Plan

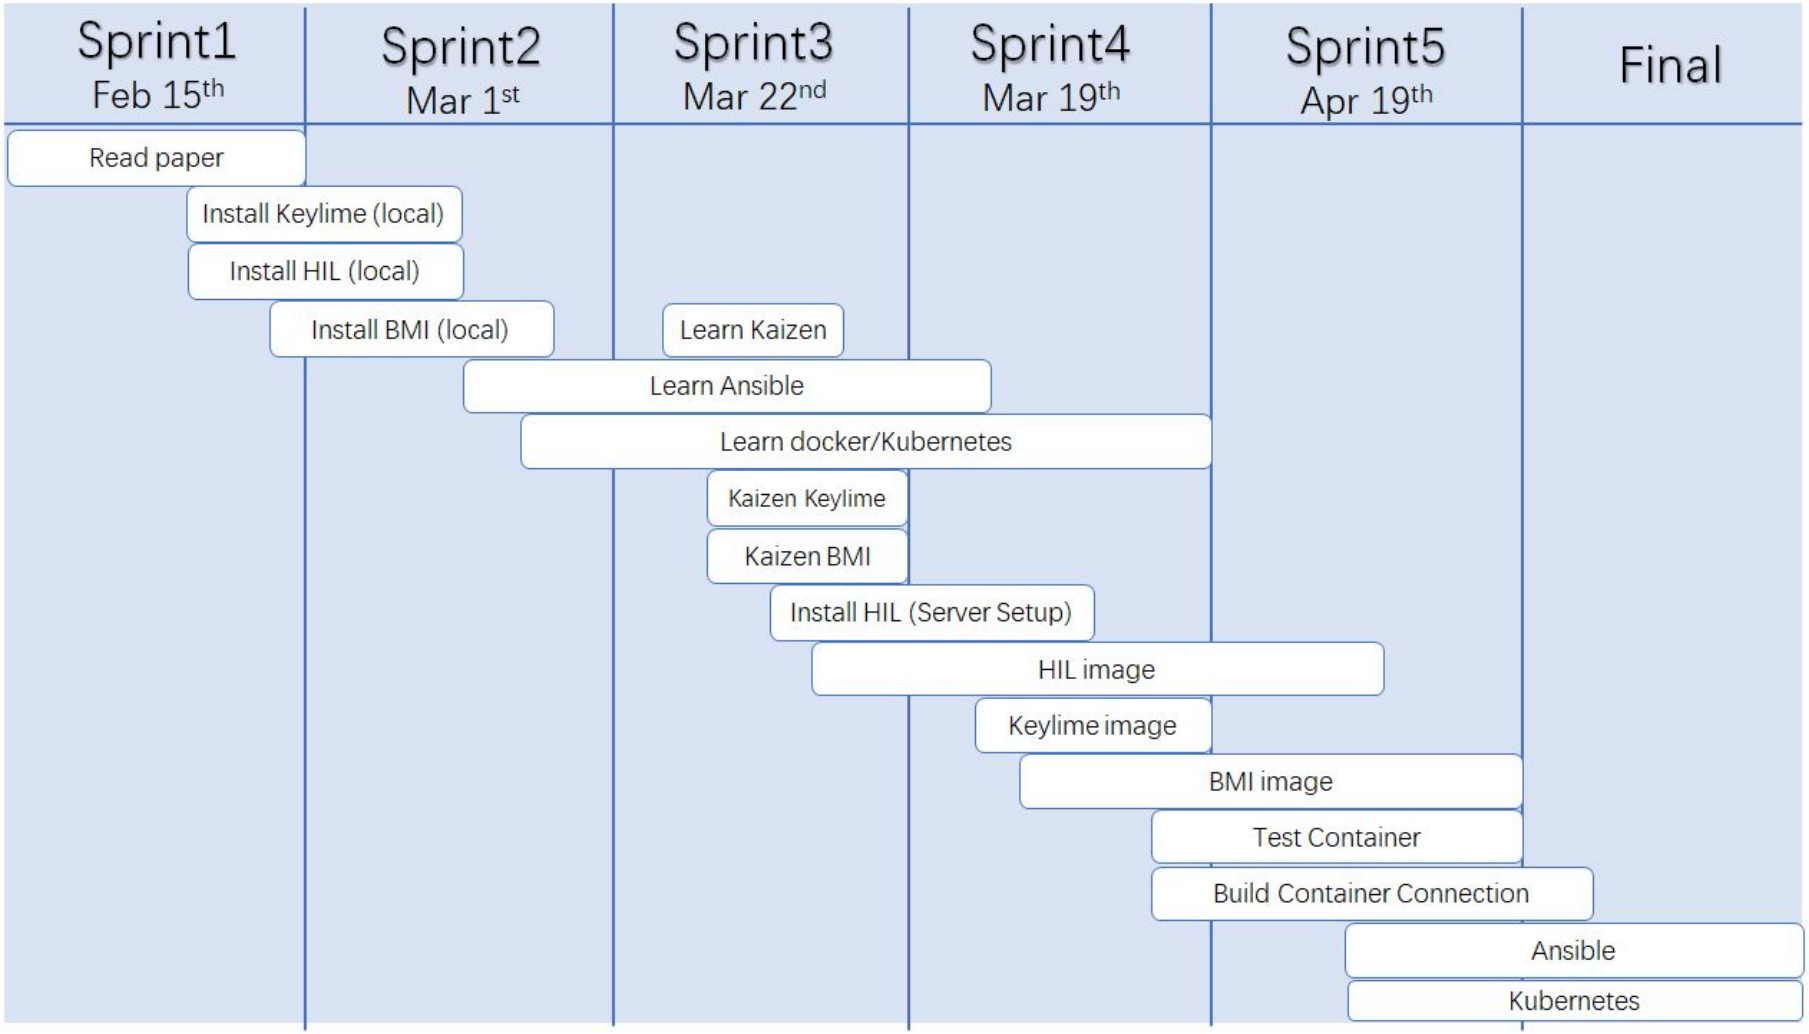

Last Sprint Report

- Finished building container image of HIL and BMI, now these two are under testing
- Learning and installing Kubernetes (In Progress)
- Found a solution to deploy containers automatically using Ansible as deployment method
- Built necessary connection between components (BMI & HIL, BMI & Ceph, HIL Servers and etc)
- Tested Keylime running on multiple VMs
- Tested HIL and BMI containers running on multiple VMs

BU

Burndown Chart

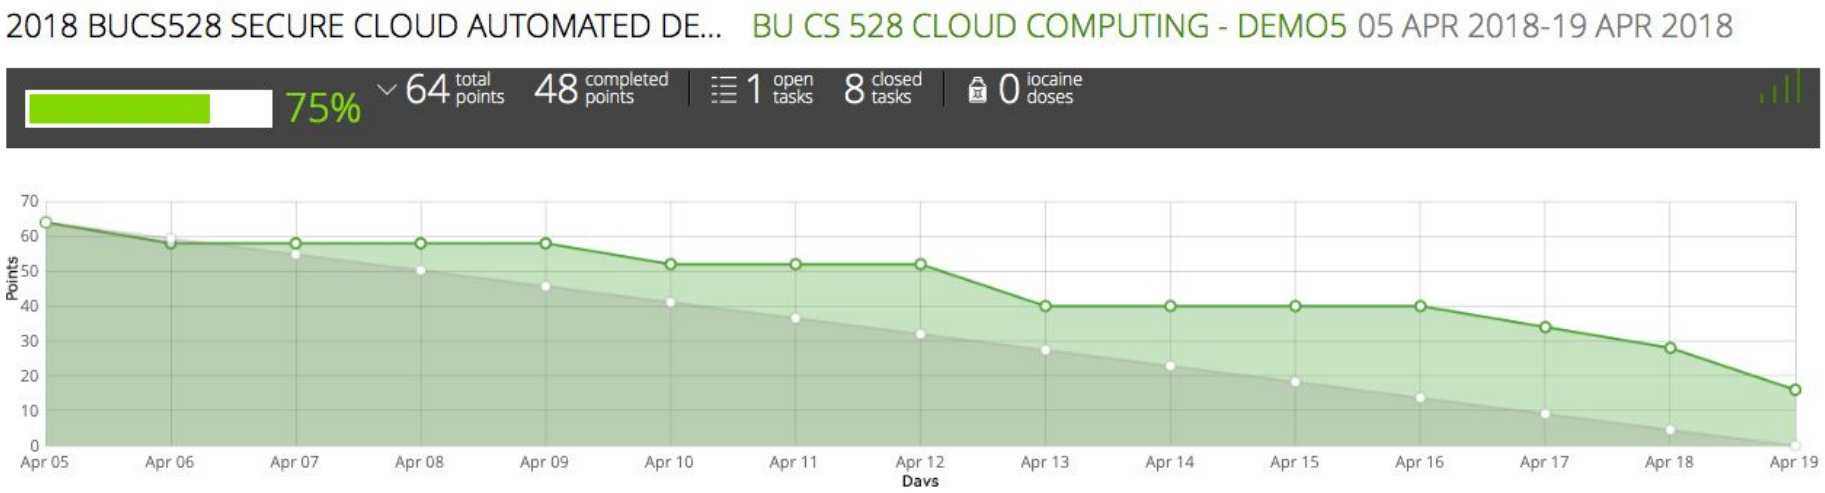

Project Progress (BMI)

- Finished deploying ceph server on Kaizen
  - One-node ceph for testing purpose
- Finished deploying BMI container on Kaizen
  - Require further configuration
  - Finished configuring ISCSI, DHCP, Sqlite3
  - Problem: Local or
- Established connection between BMI and ceph
- Have settled most of the configuration problem
- 90%

BU

Project Progress (HIL)

Solution for HIL servers is using two containers to fulfill the requirements that HIL needs to run its server.

- PostgreSql container as hil database server
- Apache/httpd container as hil wsgi apache server and network server

Two containers are built separately under the same LAN environment as HIL server

Project Progress (Ansible)

Basic concept

- **Ansible** is software that automates software provisioning, configuration management, and application deployment.
- configuration management: Mange software on top of hardware.
- Features:
  - Agentless
  - Build on top of Python
  - Use ssh for secure connection
  - Push based architecture
  - Simply
- Write playbook ---> Run playbook

Host inventory

- Contains list of hosts, grouped together.
- Default location is

  /etc/ansible/hosts

Installation

- sudo pip install ansible
- On RedHat/CentOS systems, python-pip and ansible are available via the EPEL repository
- rpm -ivh http://dl.fedoraproject.org/pub/epel/7/x86_64/\

Project Progress (Ansible)

Ping Pong between VMs

- Success ping between to VM.

```
[root@vm007 ~]# ping 10.0.0.9
PING 10.0.0.9 (10.0.0.9) 56(84) bytes of data.
64 bytes from 10.0.0.9: icmp_seq=1 ttl=64 time=2.05 ms
64 bytes from 10.0.0.9: icmp_seq=2 ttl=64 time=0.624 ms
64 bytes from 10.0.0.9: icmp_seq=3 ttl=64 time=0.538 ms
64 bytes from 10.0.0.9: icmp_seq=4 ttl=64 time=0.662 ms
64 bytes from 10.0.0.9: icmp_seq=5 ttl=64 time=0.478 ms
^C
--- 10.0.0.9 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4001ms
rtt min/avg/max/mdev = 0.478/0.872/2.059/0.597 ms
```

- Still unable to ping using ansible command

```
[root@vm007 ~]# ansible -m ping 10.0.0.9
10.0.0.9 | UNREACHABLE! => {
    "changed": false,
    "msg": "Failed to connect to the host via ssh: Permission denied (publickey,
gssapi-keyex,gssapi-with-mic).\r\n",
    "unreachable": true
}
```

Ansible playbook

- Written in YAML

```
---
 - hosts: vm007
   user: root
   vars:
    motd_welcome: 'welcome to centos007\n'
   tasks:
    - name: sample motd
      copy:
      dest: /etc/motd
      content: "{{motd_welcome}}"
```

- No syntax error. Run failed due to previous reason

```
PLAY [vm007] ***************************************************
skipping: no hosts matched

PLAY RECAP *****************************************************
```

Project Progress (Kubernetes) Challenges:

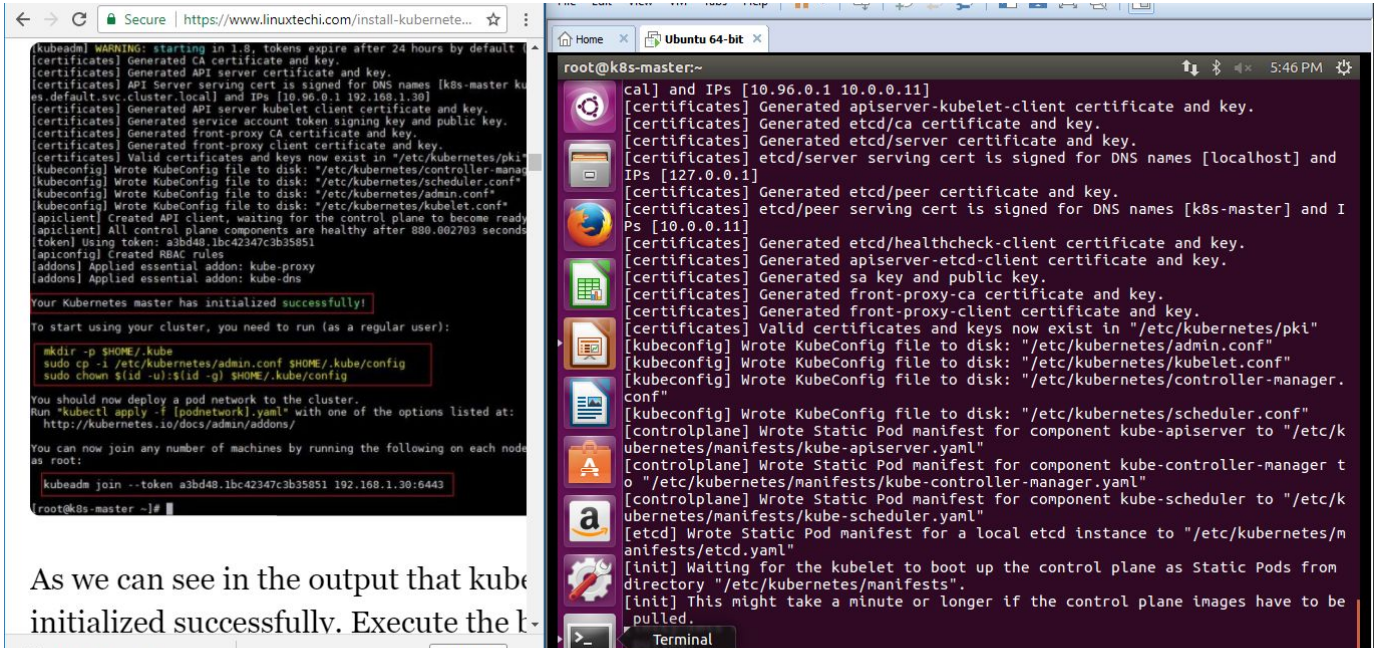

YAML file for Keylime Pod- Volume sharing between two containers:

```yaml
apiVersion: v1
kind: Pod
metadata:
  name: two-containers
spec:

  restartPolicy: Never

  volumes:
  - name: shared-data
    emptyDir: {}

  containers:

  - name: keylime
    image: docker1
    volumeMounts:
    - name: shared-data
      mountPath: /usr/share/docker/html

  - name: run.sh
    image: docker2
    volumeMounts:
    - name: shared-data
      mountPath: /pod-data
    command: ["/bin/sh"]
    args: ["-c", "echo Hello from the docker container > /pod-data/index.html"]



kubectl create -f https://k8s.io/docs/tasks/access-application-cluster/two-container-pod.yaml
kubectl get pod two-containers --output=yaml
```

Demo

User Scenario

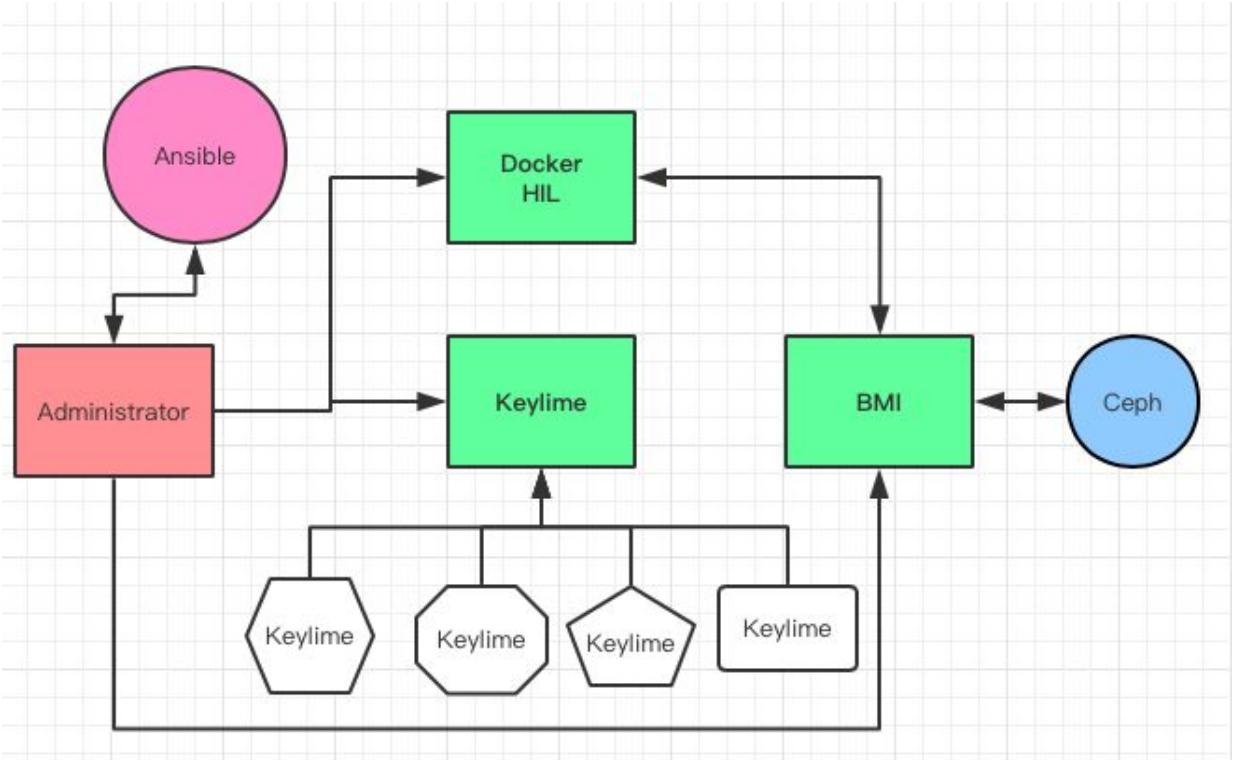

Responsibilities for next sprint

- Test and maintain container image for deploy
- An automated ansible script for installation of docker, HIL server, BMI server, Keylime server onto each VM from an admin VM for testing automated deploy
- Automate configuration between each component on Ansible Playbook based on user scenario

BU

Thank you

Question ?